Artificial intelligence and land taxation

A discussion of future scenarios

Gilles Duranton
Brief essay · Draft for discussion · June 2026

Abstract. Artificial intelligence (AI) will be remembered less for what it automates than for how it redistributes the surplus across the factors of production. Under the logic of spatial equilibrium, a sizeable share of any productivity gain ends up capitalized into the value of land. This revives, with renewed force, the old Georgist question: if land rent is a surplus the owner did not produce, is it not the tax base par excellence? The idea that taxing the rents of fixed factors becomes more attractive as AI erodes the labor base is already in the literature —from Korinek and Stiglitz (2017/2019) to Korinek and Lockwood (2026). My contribution is twofold: to show why the surplus capitalizes *preferentially* into land rather than neutrally across fixed factors; and to argue that AI itself, by rendering the separate assessment of land trivial, turns the Georgist tax from a theoretical elegance into a workable instrument. I discuss three scenarios —hyper-agglomeration, the great dispersion, and the assessment revolution— and argue that, whatever spatial pattern prevails, the case for taxing the value of land grows stronger, not weaker.

1. The starting point

It is worth beginning with what we know with some confidence. First, land —in its location component— is a factor in fixed supply: no more central land is produced when its price rises. Second, in spatial equilibrium households and firms are roughly indifferent across locations once wages, prices and amenities are netted out; differences in productivity or quality of life across places do not vanish but are *capitalized*. The corollary is uncomfortable but robust: when an innovation raises the productivity of a place, much of the gain accrues not to those who work or invest there, but to those who own the land.

These two observations are enough to suspect that AI and land taxation are more intertwined than the public conversation —today dominated by job automation— suggests. If AI is, as many anticipate, a general-purpose technology, its long-run effect on the relative prices of factors will be first-order. And among those factors, land occupies a peculiar position: it is the only one whose supply does not respond, and therefore the only one whose taxation introduces, in principle, no efficiency distortions. Henry George intuited this a century and a half ago; what is new is the context.

2. The capitalization argument

Suppose AI raises labor productivity in the occupations that concentrate in cities —professional services, finance, design, research. In the short run, this lifts wages and profits. But if productive locations are scarce —and they are, because well-located land is scarce— competition for access to them pushes rents up. Part of what AI "gives away" in productivity seeps, via real-estate prices, toward landowners. It is the same mechanism that turned the productivity gains of technology hubs into prohibitive rents: innovation is socialized as progress and privatized as land rent.

This point has a distributive consequence worth underscoring. Much of the contemporary concern about AI is that it shifts rents from labor to capital —to those who own the models, the data and the compute. That is a legitimate fear. But capital, unlike land, is reproducible: if its return rises, more capital is accumulated and the return tends to normalize. Land does not. So, over long horizons, a disproportionate fraction of the AI surplus may end up not in capital but in land, the only factor that cannot multiply in response. If so, a fiscal policy that watches only corporate capital would be guarding the wrong door.

3. Who has said this before? Situating the argument

It is worth being explicit about which part of this is new and which part takes up a conversation already under way. The idea that AI revives the case for taxing the rents of fixed factors is not mine. Its most influential formulation is that of **Korinek and Stiglitz** (2017/2019): as AI erodes labor's share of income —and with it the labor tax base— taxing the rents of factors that cannot be reproduced (unimproved land, spectrum, unique datasets) becomes at once more necessary and more efficient. Their technical point is forceful: the Atkinson–Stiglitz result, which in normal economies sidelines the land tax, loses its grip precisely when labor ceases to be a viable tax base; taxing pure rent introduces no efficiency cost, while maximum labor revenue as a share of output tends to zero as capital's share tends to one.

This line has consolidated. **Korinek and Lockwood** (2026), in their public-finance framework for the age of AI, again place land among the natural bases —alongside the excess profits of concentration and other rents— and discuss alternatives such as sovereign funds investing in AI firms or capital levies. Recent work by the **IMF** (2025) on AI adoption and inequality, and proposals such as the "algorithmic wealth tax" (2026), share the underlying diagnosis: the surplus migrates toward capital, data and rents, and the fiscal system must follow it. On the formal side, the modeling of the land value tax is also enjoying a revival of its own —the tradition of the **Henry George theorem** (Arnott and Stiglitz, 1979), by which, in the optimally sized city, aggregate land rents suffice to finance public goods, which is directly pertinent if AI is conceived as a productivity shock of a quasi-public character.

Perhaps the most telling thing is that the argument no longer lives in academia alone: it is voiced, almost word for word, by the very builders of AI. In *Moore's Law for Everything* (2021), **Sam Altman** argues that AI will make almost everything cheaper —because the dominant cost along the chain is labor— so that the two great future sources of wealth will be the companies that use AI and land, the only asset in fixed supply. His proposal, the "American Equity Fund," taxes companies and land at 2.5% per year and distributes the proceeds. The decisive thing, for our purposes, is the footnote to his reasoning: Altman acknowledges that *the theoretically optimal system would be to tax the value of land alone, not the improvements*, and that he proposes taxing both at a lower rate only because assessing land separately seems to him, "in practice," too difficult. That is: one of the principal architects of AI subscribes to the Georgist ideal and names, as his sole objection, exactly the bottleneck that this essay argues AI dissolves. No surprise, then, that Altman is the lead investor in ValueBase, Lars Doucet's land-assessment company. The circle closes: AI capital funds the very tool that renders trivial the objection Altman himself invokes. Others from the same world —Vinod Khosla, or OpenAI's fiscal-policy document, which proposes shifting the burden from labor to capital and to sustained AI-driven returns— push in the same direction, though without reaching land with Altman's clarity.

Where, then, does this essay's own contribution lie? In two shifts. *First*, the public-finance literature treats land as *one more* item on a list of fixed factors whose rents are worth taxing once the labor base is exhausted. The spatial argument is stronger: by spatial equilibrium, AI's productivity gains are not distributed at random across fixed factors but *capitalize preferentially into land*, the only strictly non-reproducible factor and therefore the residual claimant of last resort. Land is not an item on the list: it tends to be the sink where the surplus finally accumulates. *Second*, and rarely underscored in this literature, AI changes not only *how much* there is to tax but *how feasible* it is to do so: it dissolves the administrative bottleneck that for a century was the decisive objection to the Georgist tax. It is no accident that the contemporary Georgist revival —Lars Doucet, his book *Land Is a Big Deal*, and the land-assessment company ValueBase— combines Georgist advocacy with machine learning, and that its lead investor comes, with some irony, from the world of AI itself. The technique that concentrates the rent is the same one that makes it measurable.

4. Three scenarios

The direction of the effect on rents is relatively clear; its geography, by contrast, is not. Here the empirical economist's honesty is in order: we do not yet know how AI will reorganize space. I propose three scenarios, not as predictions but as limiting cases useful for thinking.

4.1. Hyper-agglomeration

If AI *complements* high-skill talent and dense in-person interaction —if the best outcomes arise where powerful models combine with concentrations of people who know how to use them— the superstar cities become even more dominant. The land-price gradient steepens. In this world, rents concentrate geographically and inequality across places sharpens. The land value tax appears here as the natural instrument for recapturing, on behalf of the community, a surplus that the community itself generated by agglomerating.

4.2. The great dispersion

If, on the contrary, AI *substitutes* for proximity —if it makes being in the same place dispensable for coordinating, deciding and creating— the gradient flattens. Value migrates from the center toward a more extensive periphery and toward intermediate cities once left behind. One would be tempted to conclude that there is then "no rent left to tax." That would be a mistake. The rent does not disappear: it is redistributed across the territory and, frequently, becomes *less visible* and therefore easier to capture regressively through speculation. A tax system anchored in the value of land adapts automatically to that redistribution; one anchored in transactions or in built improvements does not.

4.3. The assessment revolution

The third scenario is the least discussed and, I suspect, the most consequential. For a century, the practical objection to the land value tax was not theoretical but administrative: separating the value of land from the value of improvements is difficult, costly and contestable. That was the bottleneck. AI dissolves it. Satellite imagery, transaction records, mass-appraisal models and machine learning now make it possible to estimate land-value surfaces with a resolution, frequency and transparency unthinkable a decade ago. The same technology that reorders the rent provides, almost as a byproduct, the instrument to measure it. For the first time, the Georgist tax ceases to be a theoretical elegance and becomes administratively trivial.

5. Policy implications

From the three scenarios I draw one common conclusion and a couple of warnings. The conclusion: whatever spatial pattern AI imposes —concentration, dispersion or something in between— the case for taxing the value of land is reinforced. Along every path, AI generates surplus, that surplus tends to capitalize into land, and we can now measure it. A recurrent tax on the value of land —not on the building, which is produced and should be encouraged— raises revenue without distortion, discourages speculation in idle land, and returns to the community a value the community created.

First warning: the technical capacity to assess is not the same as the political capacity to tax. Landowners are, almost by definition, an organized and vocal group. AI removes the administrative excuse, but not the distributive resistance; it might even aggravate it if ownership of land becomes more concentrated. Second warning: mass algorithmic assessment shifts power to whoever controls the models and the data. An opaque tax administration, dependent on a private provider of "land prices," would reproduce within the state the very asymmetry we fear in the market. Methodological transparency and public access to assessment data are not a technical detail: they are the condition of legitimacy for the whole scheme.

6. A note of caution

I would close with the prudence the craft advises. Almost everything above rests on spatial-equilibrium intuitions and historical analogies, not on evidence about a technology that is barely beginning to deploy. We do not know the magnitude of the productivity gains, nor what fraction will capitalize into land versus wages or profits, nor at what speed. The empirical agenda is enormous and almost empty: measuring the elasticity of land prices to productivity shocks attributable to AI, distinguishing complementarity from substitution in the use of space, evaluating how algorithmic assessments behave under political pressure.

But uncertainty about magnitudes must not be confused with uncertainty about direction. The logic is clear: AI produces surplus, the surplus seeks land, and at last we have the rule to measure it. That the twenty-first century should end up remembering Henry George would be not the least of the ironies artificial intelligence has in store for us.

References

Altman, S. (2021). "Moore's Law for Everything." (Proposes the *American Equity Fund*: taxing companies and land at 2.5% per year; acknowledges that taxing land value alone would be optimal but dismisses it on grounds of assessment difficulty.)

Arnott, R. and J. E. Stiglitz (1979). "Aggregate Land Rents, Expenditure on Public Goods, and Optimal City Size," *Quarterly Journal of Economics* 93(4) (Henry George theorem).

Doucet, L. (2022). *Land Is a Big Deal*. (See also the "Does Georgism Work?" series in *Astral Codex Ten*, and the land-assessment company ValueBase.)

Duranton, G. and D. Puga (2004). "Micro-foundations of Urban Agglomeration Economies," *Handbook of Regional and Urban Economics*, vol. 4.

Duranton, G. and D. Puga (2020). "The Economics of Urban Density," *Journal of Economic Perspectives* 34(3).

George, H. (1879). *Progress and Poverty*.

International Monetary Fund (2025). "AI Adoption and Inequality," *IMF Working Paper* 2025/068.

Korinek, A. and J. E. Stiglitz (2019). "Artificial Intelligence and Its Implications for Income Distribution and Unemployment," in *The Economics of Artificial Intelligence: An Agenda* (NBER); earlier version NBER Working Paper 24174 (2017).

Korinek, A. and L. Lockwood (2026). "Public Finance in the Age of AI: A Primer," NBER Working Paper 34873 / Brookings.

A note on the literature. The argument that taxing the rents of fixed factors —land among them— becomes more attractive as AI erodes the labor tax base belongs, above all, to Korinek and Stiglitz (2017/2019) and to the public-finance agenda that follows from it (Korinek and Lockwood, 2026). What this essay adds is, first, the spatial emphasis —why the surplus capitalizes *preferentially* into land rather than spreading neutrally across fixed factors— and, second, the feasibility argument: AI dissolves the historical administrative obstacle to the separate assessment of land, moving the Georgist tax from the terrain of theory to that of applicable policy.

Note: this text is a brief discussion essay written in the style and analytical tradition of urban economics. It is not a peer-reviewed publication nor an official statement by the author cited.